


GEMM optimization ladder

each rung is one attributable step; percent is of cuBLAS at the same precision



- FP32 kernel (percent of cuBLAS SGEMM, about 23 TFLOP/s)
- tensor core kernel (percent of cuBLAS FP16 or BF16, about 65 TFLOP/s)
- top kernel